

911 WRPS Calls Data Final Project_Konsur

The Final Project the description of the dataset, the main 911 Waterloo Regional Police Service (WRPS) data sources in (*.csv) format and represents with more than 170806 incidents with 23 feature listed below. All properties of incidents are continuous variables, and Final_Priority (Severity) is an ordinal, dependent variable
http://www.wrps.on.ca/sites/default/files/WRPSOccurrenceData_Year_2014_October_15.csv
(http://www.wrps.on.ca/sites/default/files/WRPSOccurrenceData_Year_2014_October_15.csv)

The data contains the following fields:

- Occurrence Number
- Geographic Location [long, lat]
- Nearest Intersection Location
- Patrol Division
- Patrol Zone
- Municipality
- Reported [Date and Time]
- Initial Call Type (REMOVED)
- Initial Call Type Description
- Final Call Type
- Final Call Type Description
- Initial Priority (Removed)
- Final Priority [1-9]
- Disposition (Removed)
- Dispatch [Date and Time]
- Arrival [Date and Time]
- Cleared [Date and Time]
- Call Dispatch Delay = Dispath - Reported
- Call Travel Time = Arrival - Dispatch
- Call On-Scene Time = Cleared - Arrived
- Call Response Time = Arrival - Reported
- Call Service Time = Cleared - Dispatch
- Total Unit Service Time = ?????

Note: -- installing nn on my computer (from conda prompt): conda update scikit-learn followed by: pip install scikit-neuralnetwork

Data and Setup

Import numpy and pandas

```
In [187]: %matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
```

```
In [188]: from multiprocessing import Process, freeze_support
from pylab import rcParams
from scipy import stats
from sklearn import tree
from sklearn import datasets, linear_model, preprocessing
from sklearn import cross_validation
from sklearn.cross_validation import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, ExtraTreesClassifier
from sklearn.grid_search import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn import svm
from sklearn.neural_network import MLPClassifier
```

Read in the csv file as a dataframe called df


```
In [189]: path='C:\\Users\\cck023.DS\\OneDrive\\MSCA\\data\\WRPSOccurrenceData_Year_2015_
_November3_2016.csv'
```

```
In [190]: df=pd.read_csv(path)
df.head(3)
```

Out[190]:

	Occurrence Number	Geographic Location	Nearest Intersection Location	Patrol Division	Patrol Zone	Municipality	Reported Date and Time
0	WA15168339	537724.704700 ,4816634.248000	KING ST N / NORTHFIELD DR E / NORTHFIELD DR W	WN	WN6	NaN	2015-08-08 12:00:00
1	WA15274061	552814.080200 ,4804681.021100	MONTROSE ST N / RAILWAY ST	WS	WS6	CAM	2015-12-24 12:15:17
2	WA15240889	543979.878500 ,4805604.983000	BEASLEY DR / HOMER WATSON BLVD	WC	WC7	KIT	2015-11-09 13:19:15

3 rows × 23 columns



Check the info() of the df

```
In [191]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 280141 entries, 0 to 280140
Data columns (total 23 columns):
Occurrence Number      280141 non-null object
Geographic Location     275947 non-null object
Nearest Intersection Location 275947 non-null object
Patrol Division         279233 non-null object
Patrol Zone             275594 non-null object
Municipality            185449 non-null object
Reported Date and Time  280141 non-null object
Initial Call Type       186435 non-null float64
Initial Call Type Description 186435 non-null object
Final Call Type         280138 non-null float64
Final Call Type Description 280138 non-null object
Initial Priority         186435 non-null float64
Final Priority           280139 non-null float64
Disposition              186452 non-null object
Dispatch Date and Time  186468 non-null object
Arrival Date and Time   176577 non-null object
Cleared Date and Time   186417 non-null object
Call Dispatch Delay     186468 non-null float64
Call Travel Time        176574 non-null float64
Call On-Scene Time      176537 non-null float64
Call Response Time      176577 non-null float64
Call Service Time       186413 non-null float64
Total Unit Service Time 186455 non-null float64
dtypes: float64(10), object(13)
memory usage: 49.2+ MB
```

Check the head of df

```
In [192]: df1=df.dropna()
```

```
In [193]: df1.shape
```


```
Out[193]: (170806, 23)
```

```
In [194]: df1.reset_index(drop=True)
df1.head(5)
```

Out[194]:

	Occurrence Number	Geographic Location	Nearest Intersection Location	Patrol Division	Patrol Zone	Municipality	Reported Date and Time
1	WA15274061	552814.080200 ,4804681.021100	MONTROSE ST N / RAILWAY ST	WS	WS6	CAM	2015-12-24 12:15:17
2	WA15240889	543979.878500 ,4805604.983000	BEASLEY DR / HOMER WATSON BLVD	WC	WC7	KIT	2015-11-09 13:19:15
3	WA15171651	556059.710500 ,4802351.015300	COULSON ST / ELMWOOD AVE	WS	WS4	CAM	2015-08-13 07:38:07
4	WA15200283	541176.033600 ,4810854.312200	CHARLES ST W / ONTARIO ST S	WC	WC1	KIT	2015-09-18 14:30:35
6	WA15162132	538320.095900 ,4814442.414200	COLUMBIA ST E / COLUMBIA ST W / KING ST N	W3	324	WAT	2015-07-31 07:01:43

5 rows × 23 columns



```
In [195]: df1.shape
```

Out[195]: (170806, 23)

```
In [196]: df1 = df1.drop('Disposition', 1)
df1 = df1.drop('Initial Priority', 1)
```

```
In [197]: df1 = df1.drop('Initial Call Type', 1)
```

```
In [198]: df1 = df1.drop('Initial Call Type Description',1)
```

```
In [199]: df1.shape
```

Out[199]: (170806, 19)

In [200]: df1.head(5)

Out[200]:

	Occurrence Number	Geographic Location	Nearest Intersection Location	Patrol Division	Patrol Zone	Municipality	Reported Date and Time
1	WA15274061	552814.080200 ,4804681.021100	MONTROSE ST N / RAILWAY ST	WS	WS6	CAM	2015-12-24 12:15:17
2	WA15240889	543979.878500 ,4805604.983000	BEASLEY DR / HOMER WATSON BLVD	WC	WC7	KIT	2015-11-09 13:19:15
3	WA15171651	556059.710500 ,4802351.015300	COULSON ST / ELMWOOD AVE	WS	WS4	CAM	2015-08-13 07:38:07
4	WA15200283	541176.033600 ,4810854.312200	CHARLES ST W / ONTARIO ST S	WC	WC1	KIT	2015-09-18 14:30:35
6	WA15162132	538320.095900 ,4814442.414200	COLUMBIA ST E / COLUMBIA ST W / KING ST N	W3	324	WAT	2015-07-31 07:01:43

In [201]: df1.columns = [c.replace(' ', '_') for c in df1.columns]

In [202]: df1.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 170806 entries, 1 to 280140
Data columns (total 19 columns):
Occurrence_Number      170806 non-null object
Geographic_Location     170806 non-null object
Nearest_Intersection_Location 170806 non-null object
Patrol_Division         170806 non-null object
Patrol_Zone             170806 non-null object
Municipality            170806 non-null object
Reported_Date_and_Time  170806 non-null object
Final_Call_Type         170806 non-null float64
Final_Call_Type_Description 170806 non-null object
Final_Priority          170806 non-null float64
_Dispatch_Date_and_Time 170806 non-null object
Arrival_Date_and_Time   170806 non-null object
Cleared_Date_and_Time   170806 non-null object
Call_Dispatch_Delay     170806 non-null float64
Call_Travel_Time        170806 non-null float64
Call_On-Scene_Time      170806 non-null float64
Call_Response_Time      170806 non-null float64
Call_Service_Time       170806 non-null float64
Total_Unit_Service_Time 170806 non-null float64
dtypes: float64(8), object(11)
memory usage: 26.1+ MB
```

In [203]: df1['Reason'] = df1['Final_Call_Type_Description'].apply(lambda Final_Call_Type_Description: Final_Call_Type_Description.split('-')[1])

```
In [204]: df1['Reason'].value_counts()
```


Out[204]:

VEHICLE STOP	41624
COMPASSIONATE TO LOCATE	12589
BYLAW COMPLAINT	8460
THEFT UNDER \$5000	6132
SELECTIVE TRAFFIC ENFORCEMENT PROGRAM (STEP)	5957
DOMESTIC DISPUTE	5610
ADMINISTRATIVE/ROUTINE DETAIL	5599
PROACTIVE INITIATIVE/PROJECT	5583
UNWANTED PERSON	5525
MVC PROP. DAMAGE	4874
DRIVING COMPLAINT	4591
PERSON STOP	4067
ALARM	3465
INJURED/SICK PERSON	3457
DISPUTE	3425
DOMESTIC OTHER	2691
ARREST	2470
SUSPICIOUS PERSON	2350
PROP. (LOST AND FOUND)	2294
DISTURBANCE	2289
MENTALLY ILL	2155
BREAK & ENTER	2098
MISSING PERSON	2034
MVC PERSONAL INJURY	1683
FRAUD	1633
DRUGS	1606
ASSIST OTHER SERVICE	1591
PROPERTY DAMAGE	1586
ATTEMPT SUICIDE	1475
PAID DUTY	1329
...	
SEX OFFENCE	341
CHILD CUSTODY AND ACCESS	318
WORKPLACE ACCIDENT	233
ROBBERY	220
GRAFFITI	183
VEHICLE / PLATE SEIZURE	149
ALERT	140
SEARCH WARRANT	138
CRIMINAL HARASSMENT/STALKING	137
COUNTERFEIT MONEY	121
THEFT OVER \$5000	96
INDECENT ACT	95
ADMINISTRATIVE NOTICE (9	93
ELDER ABUSE	68
PROSTITUTION	62
PORNOGRAPHY	53
SUSPECT APPREHENSION PURSUIT	50
PROWLER	41
PUBLIC MISCHIEF	36
EXTORTION	20
TECHNOLOGY/INTERNET CRIME	17
BOMB THREAT	14
MVC FATALITY	12
LABOUR DISPUTE	12
ABDUCTION	6
GAMING & BETTING	6

HOMICIDE	6
FILED 9	5
HUMAN TRAFFICKING	4
UNKNOWN CALL REQUIRING POLICE ASSISTANCE	1

Name: Reason, dtype: int64

What are the top 5 Patrol Zones for 911 calls?

```
In [205]: df1['Patrol_Zone'].value_counts().head(5)
```


```
Out[205]: WS4      10229
          WN6       9935
          WC2       9508
          WC5       9341
          WS2       9315
          Name: Patrol_Zone, dtype: int64
```

```
In [206]: df1['Patrol_Division'].value_counts().head(5)
```

```
Out[206]: WN      62662
          WC      59070
          WS      47492
          WR       275
          W2       232
          Name: Patrol_Division, dtype: int64
```

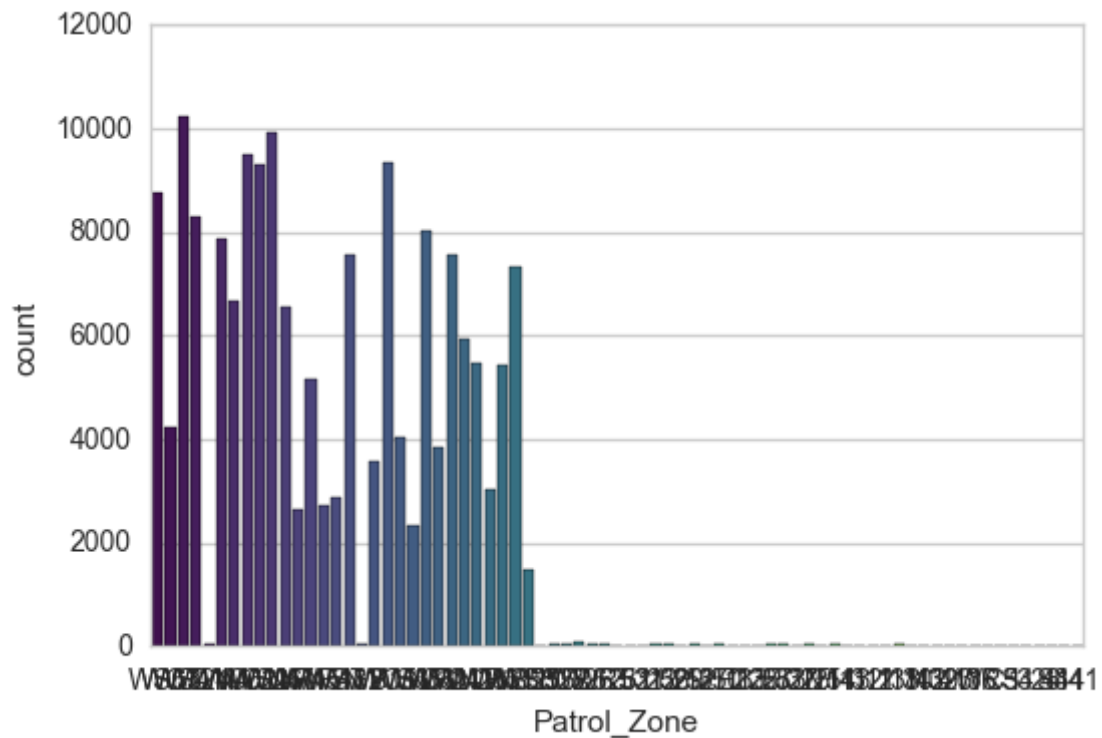
```
In [207]: sns.countplot(x='Reason',data=df1,palette='viridis')
```

```
Out[207]: <matplotlib.axes._subplots.AxesSubplot at 0x2caadbbf9e8>
```




```
In [208]: sns.countplot(x='Patrol_Zone',data=df1,palette='viridis')
```

```
Out[208]: <matplotlib.axes._subplots.AxesSubplot at 0x2caaf339a58>
```



```
In [209]: sns.countplot(x='Patrol_Division',data=df1,palette='viridis')
```

```
Out[209]: <matplotlib.axes._subplots.AxesSubplot at 0x2caa533f940>
```



Y: Recoding target variable: If 'Final Priority' >= 7, then 1, else 0.

```
In [210]: df1.shape
```

```
Out[210]: (170806, 20)
```

```
In [211]: df1.columns
```

```
Out[211]: Index(['Occurrence_Number', 'Geographic_Location',  
                'Nearest_Intersection_Location', 'Patrol_Division', 'Patrol_Zone',  
                'Municipality', 'Reported_Date_and_Time', 'Final_Call_Type',  
                'Final_Call_Type_Description', 'Final_Priority',  
                '_Dispatch_Date_and_Time', 'Arrival_Date_and_Time',  
                'Cleared_Date_and_Time', 'Call_Dispatch_Delay', 'Call_Travel_Time',  
                'Call_On-Scene_Time', 'Call_Response_Time', 'Call_Service_Time',  
                'Total_Unit_Service_Time', 'Reason'],  
               dtype='object')
```

```
In [212]: Y = df1.Final_Priority.values
```

```
In [213]: df1 = df1.drop('Final_Priority', axis=1)  
X = df1.as_matrix()  
X_cols = df1.columns[0:11]  
Y = np.asarray([1 if i>=7 else 0 for i in Y])
```

```
In [214]: X.shape
```

```
Out[214]: (170806, 19)
```

```
In [215]: Y.shape
```

```
Out[215]: (170806,)
```

```
In [216]: X[:,12:17]
```

```
Out[216]: array([[13772.0, 330.0, 278.0, 14102.0, 608.0],  
                [5.0, 1.0, 993.0, 6.0, 994.0],  
                [19.0, 270.0, 1072.0, 289.0, 1342.0],  
                ...,  
                [19403.0, 433.0, 1188.0, 19836.0, 1621.0],  
                [227.0, 717.0, 3633.0, 944.0, 4350.0],  
                [36.0, 7.0, 7564.0, 43.0, 7571.0]], dtype=object)
```

```
In [217]: X2=X[:,12:17]
```

In [218]: df1.head(5)

Out[218]:

	Occurrence_Number	Geographic_Location	Nearest_Intersection_Location	Patrol_C
1	WA15274061	552814.080200 ,4804681.021100	MONTROSE ST N / RAILWAY ST	WS
2	WA15240889	543979.878500 ,4805604.983000	BEASLEY DR / HOMER WATSON BLVD	WC
3	WA15171651	556059.710500 ,4802351.015300	COULSON ST / ELMWOOD AVE	WS
4	WA15200283	541176.033600 ,4810854.312200	CHARLES ST W / ONTARIO ST S	WC
6	WA15162132	538320.095900 ,4814442.414200	COLUMBIA ST E / COLUMBIA ST W / KING ST N	W3

In [219]: *#Creating train and test set.*
np.random.seed(2)
X_train, X_test, Y_train, Y_test = train_test_split(X2[1:1000,:], Y[1:1000], t
rain_size=0.7)

I'll standardize the features so that they are centered around 0 with a standard deviation of 1.

In [220]: std_scale = preprocessing.StandardScaler().fit(X_train)
X_train = std_scale.transform(X_train)
X_test = std_scale.transform(X_test)

C:\Users\cck023.DS\AppData\Local\Continuum\Anaconda3\lib\site-packages\sklearn
utils\validation.py:429: DataConversionWarning: Data with input dtype objec
t was converted to float64 by StandardScaler.
warnings.warn(msg, _DataConversionWarning)

Quick test of multiple models with cross validation

In [221]: clf1 = LogisticRegression()
clf2 = RandomForestClassifier()
clf3 = GaussianNB()
clf4 = ExtraTreesClassifier()

```
In [222]: for clf, label in zip([clf1, clf2, clf3, clf4],
    ['Logistic Regression', 'Random Forest', 'Naive Bayes', 'Extra Trees']):
    scores = cross_validation.cross_val_score(clf, X_train, Y_train, cv=10, scoring='accuracy')
    print("Accuracy: %0.2f (+/- %0.2f) [%s]" % (scores.mean(), scores.std(), label))
```

```
Accuracy: 0.87 (+/- 0.03) [Logistic Regression]
Accuracy: 0.92 (+/- 0.03) [Random Forest]
Accuracy: 0.54 (+/- 0.06) [Naive Bayes]
Accuracy: 0.89 (+/- 0.03) [Extra Trees]
```

```
In [223]: clf6 = GradientBoostingClassifier()
    clf7 = tree.DecisionTreeClassifier()
```

```
In [224]: for clf, label in zip([clf6, clf7],
    ['Gradient Boosting', 'Decision Tree']):
    scores = cross_validation.cross_val_score(clf, X_train, Y_train, cv=10, scoring='accuracy')
    print("Accuracy: %0.2f (+/- %0.2f) [%s]" % (scores.mean(), scores.std()*2, label))
```

```
Accuracy: 0.92 (+/- 0.07) [Gradient Boosting]
Accuracy: 0.90 (+/- 0.06) [Decision Tree]
```

```
In [225]: clf5 = SVC()
    clf5
```

```
Out[225]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
In [226]: for clf, label in zip([clf5],
    ['SVM']):
    scores = cross_validation.cross_val_score(clf, X_train, Y_train, cv=2, scoring='accuracy')
    print("Accuracy: %0.2f (+/- %0.2f) [%s]" % (scores.mean(), scores.std(), label))
```

```
Accuracy: 0.80 (+/- 0.01) [SVM]
```

Gradient Boosting, Random Forest and Extra trees performs the best, while naive bayes performs the worst. I will drop Logistic Regression and Naive Bayes and will analyze the following models: Decision Tree, Random Forest, Gradient Boosting, SVM, and Extra Trees. Decision Tree

DecisionTree Classifier with default parameters

```
In [227]: dt = tree.DecisionTreeClassifier(criterion="gini", min_samples_split=10)
    dt.fit(X_train, Y_train)
    dt_test_preds = dt.predict_proba(X_test)[: , 1]
```

```
In [228]: dt_preds = dt.predict(X_test)
print (pd.crosstab(index=Y_test, columns=dt_preds, rownames=["Actual"], colnames=["Pred"]))
```

Pred	0	1
Actual		
0	175	16
1	14	95

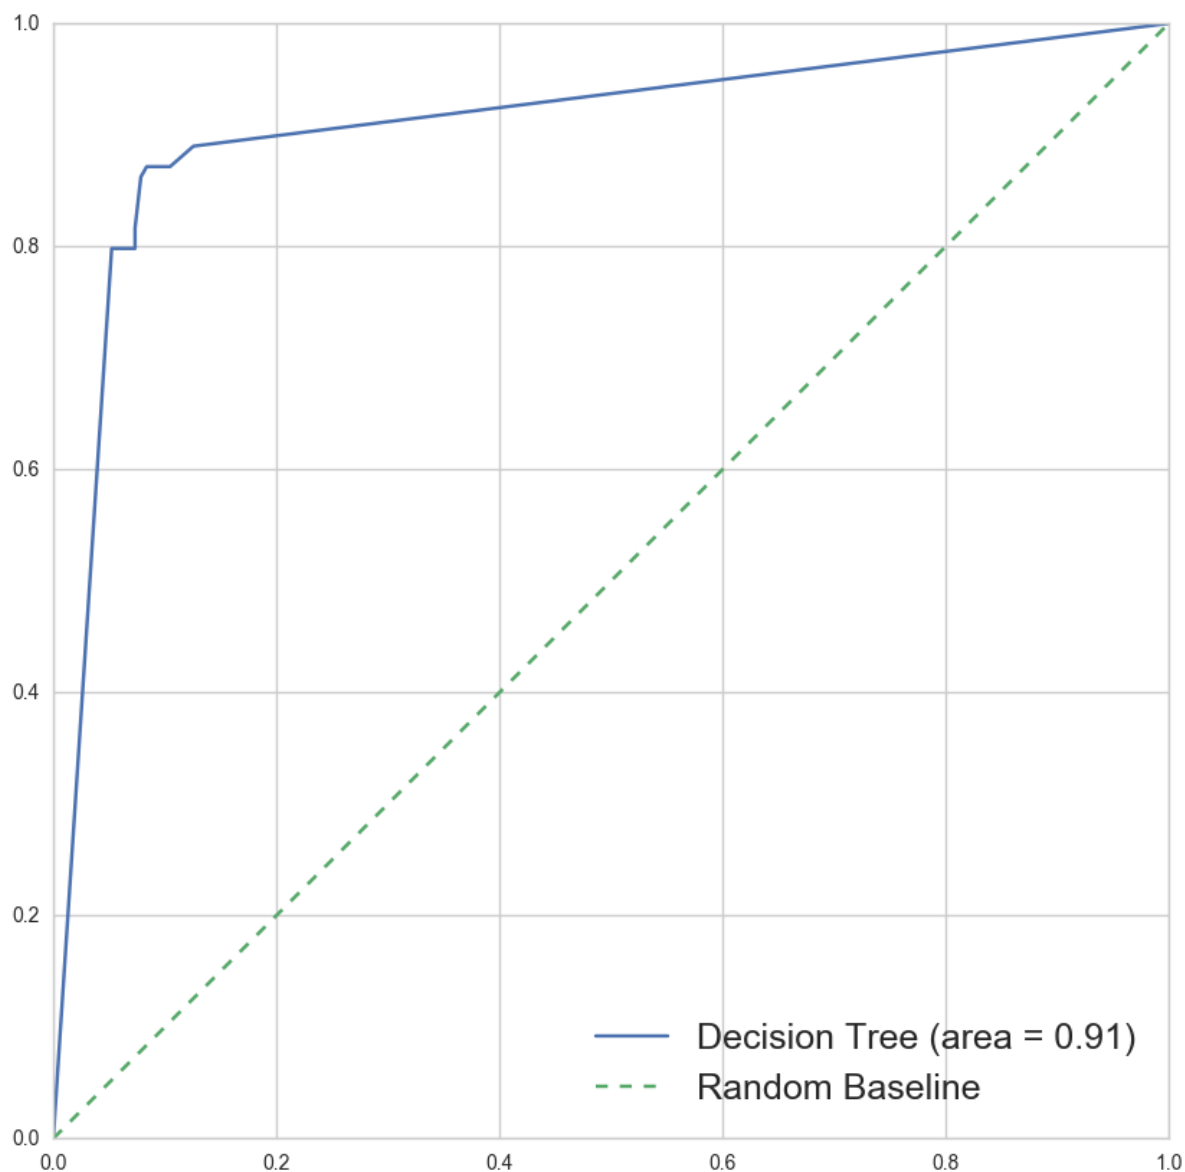
```
In [229]: dt_roc = roc_auc_score(Y_test, dt_test_preds)
dt_roc
```

```
Out[229]: 0.90945770690234884
```

```
In [230]: fpr_decision_tree, tpr_decision_tree, thresholds_decision_tree = roc_curve(Y_test, dt_test_preds)
```



```
In [231]: fpr_rand = tpr_rand = np.linspace(0, 1, 10)
rcParams['figure.figsize'] = 10, 10
plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %
0.2f)' % dt_roc)
plt.plot(fpr_rand, tpr_rand, linestyle='--', label='Random Baseline')
plt.legend(loc='lower right', prop={'size':18})
plt.show()
```



Random Forest Random Forest Classifier with default parameters

```
In [232]: rf = RandomForestClassifier()
rf.fit(X_train, Y_train)
rf_test_preds = rf.predict_proba(X_test)[: , 1]
```

```
In [233]: rf_preds = rf.predict(X_test)
print (pd.crosstab(index=Y_test, columns=rf_preds, rownames=["Actual"], colnames=["Pred"]))
```


Pred	0	1
Actual		
0	177	14
1	9	100

```
In [234]: rf_roc = roc_auc_score(Y_test, rf_test_preds)
roc_auc_score(Y_test, rf_test_preds)
```

```
Out[234]: 0.94788414429127243
```


```
In [235]: fpr_rf, tpr_rf, thresholds_rf = roc_curve(Y_test, rf_test_preds)
```

```
In [236]: plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %  
0.2f)' % dt_roc)  
plt.plot(fpr_rf, tpr_rf, label='Random Forest (area = %0.2f)' % rf_roc)  
plt.plot(fpr_rand, tpr_rand, linestyle='--', label='Random Baseline')  
plt.legend(loc='lower right', prop={'size':18})  
plt.show()
```



*Random Forest Classifier feature importance


```
In [237]: rf_importances = rf.feature_importances_
sorted_indices = np.argsort(rf_importances)
y_pos = range(len(rf_importances))
plt.barh(y_pos, rf_importances[sorted_indices], align='center')
plt.yticks(y_pos, X_cols[sorted_indices])
plt.show()
```



```
In [238]: scores = []
```

```
In [239]: for val in range(1,10):
rf = RandomForestClassifier(n_estimators = val)
validated = cross_val_score(rf, X_train, Y_train, cv=10)
scores.append(validated)
```

```
In [240]: sns.boxplot(scores)
plt.xlabel('Number of Trees')
plt.ylabel('Classification Scores')
plt.title('Classification Score for Number of Trees')
plt.show()
```



The classification score increases as the number of trees increase but it stadystate at 10.

```
In [241]: param_grid = {"n_estimators": [5,19,26],
    "max_depth": [3,4,5,6,7, None],
    "max_features": ['auto', None],
    "min_samples_split": [1,3,5,10],
    "min_samples_leaf": [1,3,5,10]}
```

```
In [242]: rf_ml = RandomForestClassifier()
```

```
In [243]: #grid_search = GridSearchCV(rf_ml, param_grid=param_grid, cv=5)
#grid_search.fit(X_train, Y_train)
#print (grid_search.best_params_)
#{'max_features': 'auto', 'min_samples_split': 1, 'n_estimators': 19, 'max_depth': None,
#'min_samples_leaf': 1}
```

```
In [244]: rf_custom = RandomForestClassifier(n_estimators=19, max_depth=None, max_features='auto',
min_samples_split=2, min_samples_leaf=2)
rf_custom.fit(X_train, Y_train)
```

```
Out[244]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_split=1e-07, min_samples_leaf=2,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=19, n_jobs=1, oob_score=False, random_state=None,
verbose=0, warm_start=False)
```

```
In [245]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_samples_leaf=1, min_samples_split=1,
min_weight_fraction_leaf=0.0, n_estimators=19, n_jobs=1,
oob_score=False, random_state=None, verbose=0,
warm_start=False)
```

```
Out[245]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_split=1e-07, min_samples_leaf=1,
min_samples_split=1, min_weight_fraction_leaf=0.0,
n_estimators=19, n_jobs=1, oob_score=False, random_state=None,
verbose=0, warm_start=False)
```

```
In [246]: rf_custom_preds = rf_custom.predict(X_test)
print (pd.crosstab(Y_test, rf_custom_preds, rownames=["Actual"], colnames=["Pred"]))
```

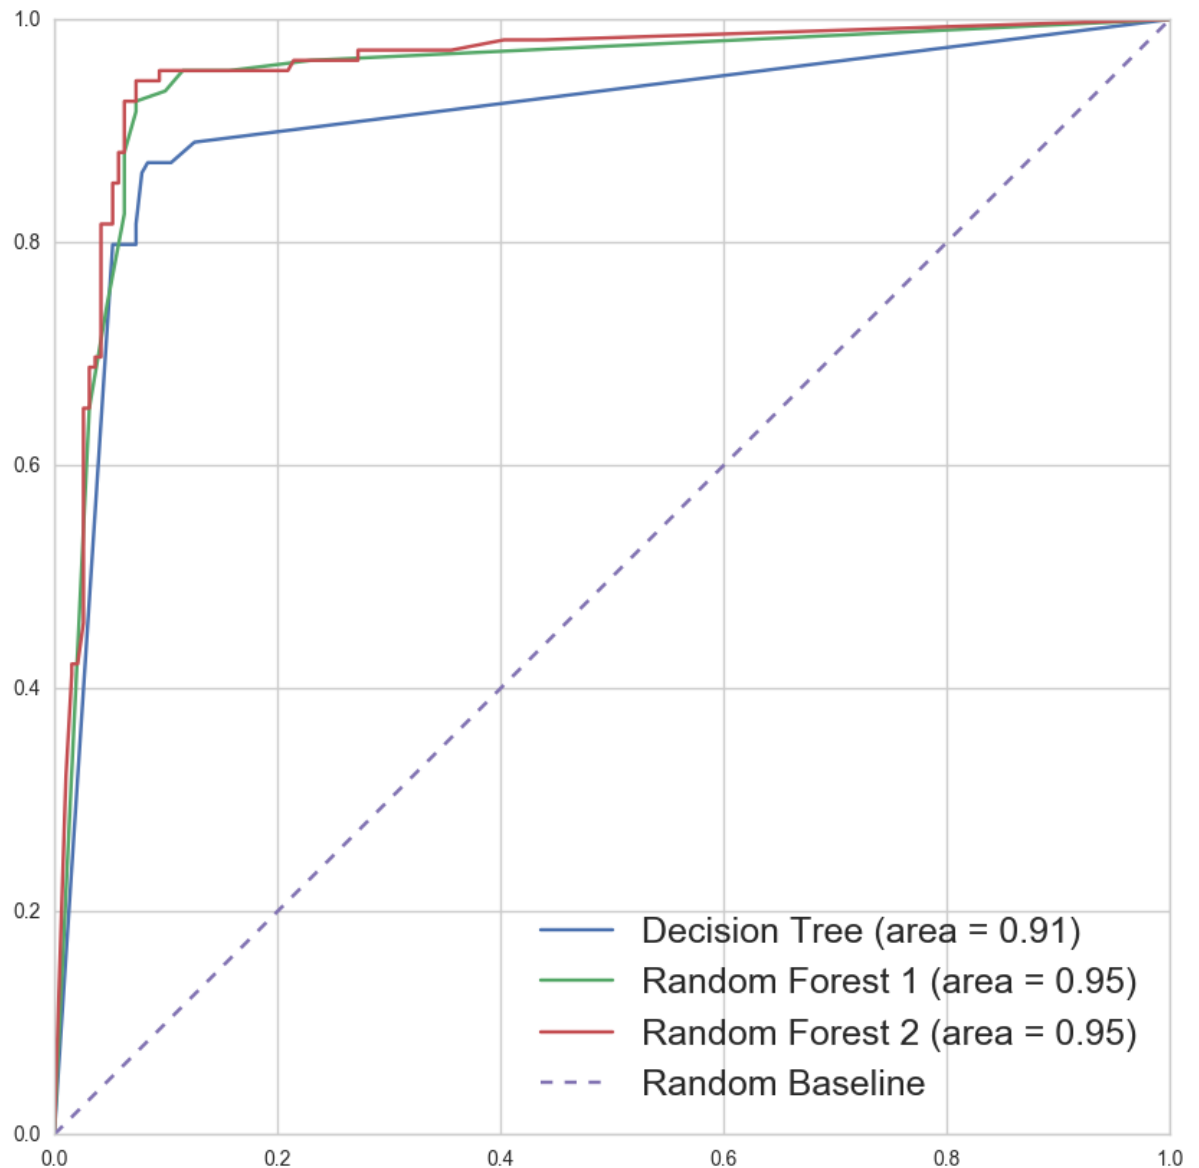
Pred	0	1
Actual		
0	177	14
1	8	101

```
In [247]: rf_custom_test_preds = rf_custom.predict_proba(X_test)[: , 1]
rf_custom_roc = roc_auc_score(Y_test, rf_custom_test_preds)
roc_auc_score(Y_test, rf_custom_test_preds)
```

```
Out[247]: 0.9548489360680148
```

```
In [248]: fpr_rf_custom, tpr_rf_custom, thresholds_rf_custom = roc_curve(Y_test, rf_custom_test_preds)
```

```
In [249]: plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %
0.2f)' % dt_roc)
plt.plot(fpr_rf, tpr_rf, label='Random Forest 1 (area = %0.2f)' % rf_roc)
plt.plot(fpr_rf_custom, tpr_rf_custom, label='Random Forest 2 (area = %0.2f)'
% rf_custom_roc)
plt.plot(fpr_rand, tpr_rand, linestyle='--', label='Random Baseline')
plt.legend(loc='lower right', prop={'size':18})
plt.show()
```



Gradient Boosting Classifier

```
In [250]: gbc = GradientBoostingClassifier(learning_rate=0.1, n_estimators=500, max_depth=5)
```

```
In [251]: gbc.fit(X_train, Y_train)
gbc_test_preds = gbc.predict_proba(X_test)[: , 1]
```

```
In [252]: gbc_test_preds = gbc.predict_proba(X_test)[: , 1]
gbc_preds = gbc.predict(X_test)
```

```
In [253]: print (pd.crosstab(index=Y_test, columns=gbc_preds, rownames=['True'], colnames=['Predicted']))
```

Predicted	0	1
True		
0	176	15
1	15	94

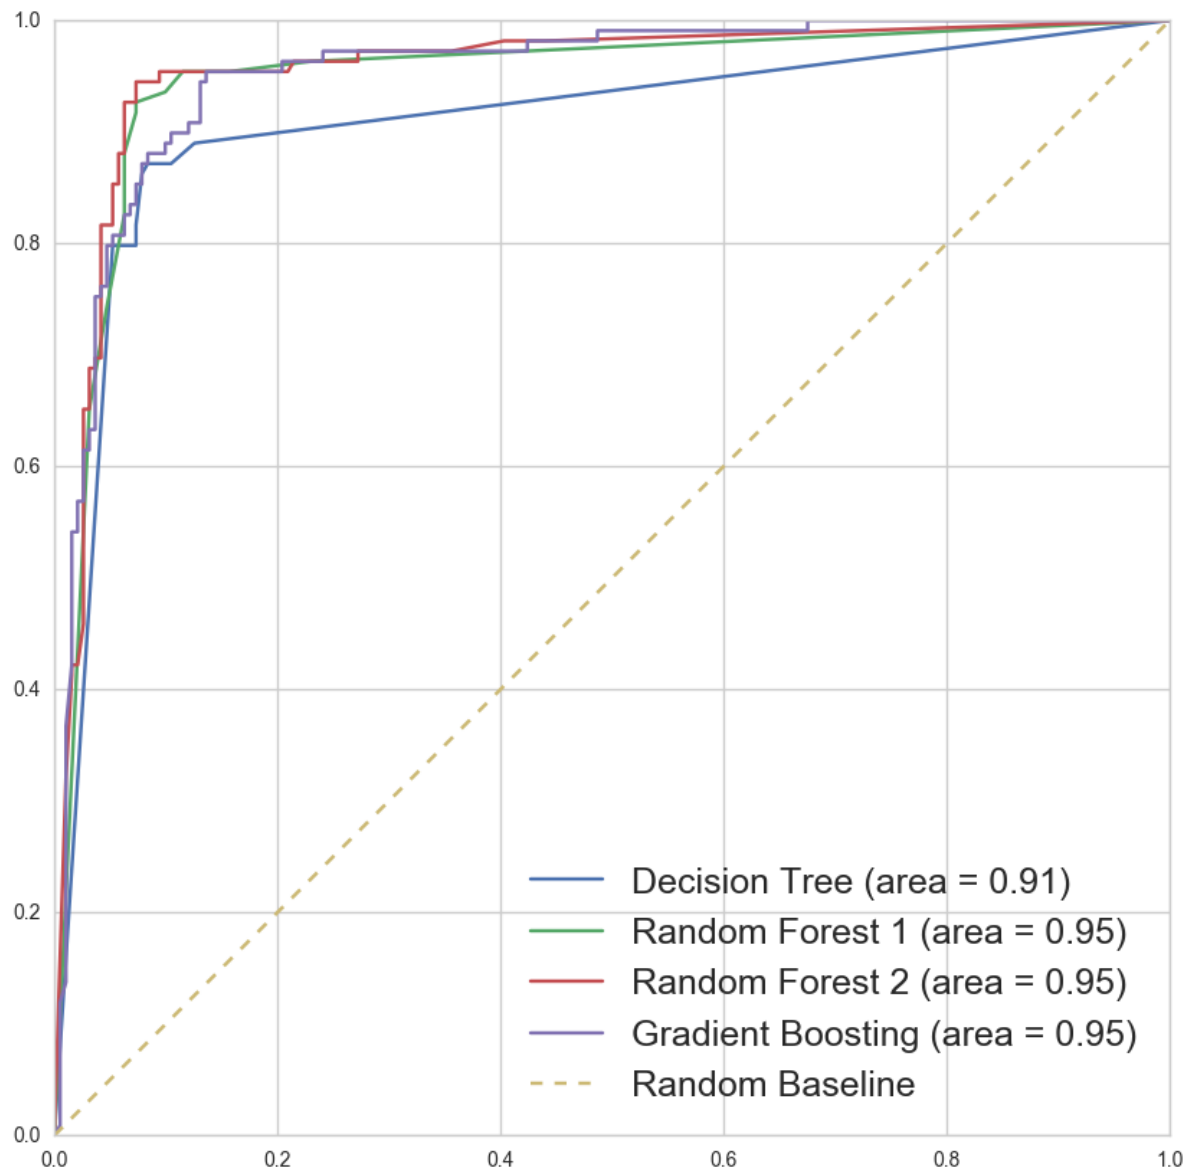
```
In [254]: gbc_roc = roc_auc_score(Y_test, gbc_test_preds)
roc_auc_score(Y_test, gbc_test_preds)
```

```
Out[254]: 0.95194293674047747
```


```
In [255]: fpr_gbc, tpr_gbc, thresholds_gbc = roc_curve(Y_test, gbc_test_preds)
```



```
In [256]: plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %  
0.2f)' % dt_roc)  
plt.plot(fpr_rf, tpr_rf, label='Random Forest 1 (area = %0.2f)' % rf_roc)  
plt.plot(fpr_rf_custom, tpr_rf_custom, label='Random Forest 2 (area = %0.2f)'  
% rf_custom_roc)  
plt.plot(fpr_gbc, tpr_gbc, label='Gradient Boosting (area = %0.2f)' % gbc_roc)  
plt.plot(fpr_rand, tpr_rand, linestyle='--', label='Random Baseline')  
plt.legend(loc='lower right', prop={'size':18})  
plt.show()
```



```
In [257]: gbc_importances = gbc.feature_importances_
sorted_indices = np.argsort(gbc_importances)
y_pos = range(len(gbc_importances))
plt.barh(y_pos, gbc_importances[sorted_indices], align='center')
plt.barh(y_pos, gbc_importances[sorted_indices], align='center')
plt.yticks(y_pos, X_cols[sorted_indices])
plt.show()
```



Support Vector Machine (SVM)

```
In [258]: clf = SVC(C=1.0, probability=True)
clf
```

```
Out[258]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
max_iter=-1, probability=True, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```
In [259]: X_train.shape
```

```
Out[259]: (699, 5)
```

```
In [260]: Y_train.shape
```

```
Out[260]: (699,)
```

```
In [261]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0, degree=3, kernel='rbf', max_iter=-1, probability=True, random_state=None, shrinking=True, tol=0.001, verbose=False)
```

```
Out[261]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0, decision_function_shape=None, degree=3, gamma='auto', kernel='rbf', max_iter=-1, probability=True, random_state=None, shrinking=True, tol=0.001, verbose=False)
```

```
In [262]: clf.fit(X_train, Y_train)
```

```
Out[262]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0, decision_function_shape=None, degree=3, gamma='auto', kernel='rbf', max_iter=-1, probability=True, random_state=None, shrinking=True, tol=0.001, verbose=False)
```

```
In [263]: svm_test_preds = clf.predict_proba(X_test)[: , 1]
svm_preds = clf.predict(X_test)
```

```
In [264]: print (pd.crosstab(index=Y_test, columns=svm_preds, rownames=['True'], columns=['Predicted']))
```

Predicted	0	1
True		
0	160	31
1	11	98


```
In [265]: svm_roc = roc_auc_score(Y_test, svm_test_preds)
roc_auc_score(Y_test, svm_test_preds)
```

```
Out[265]: 0.93839761756088191
```


```
In [266]: fpr_svm, tpr_svm, thresholds_svm = roc_curve(Y_test, svm_test_preds)
```

```
In [267]: plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %  
0.2f)' % dt_roc)
```

```
Out[267]: [<matplotlib.lines.Line2D at 0x2cab884beb8>]
```



```
In [268]: plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %
0.2f)' % dt_roc)
plt.plot(fpr_rf, tpr_rf, label='Random Forest 1 (area = %0.2f)' % rf_roc)
plt.plot(fpr_rf_custom, tpr_rf_custom, label='Random Forest 2 (area = %0.2f)'
% rf_custom_roc)
plt.plot(fpr_gbc, tpr_gbc, label='Gradient Boosting (area = %0.2f)' % gbc_roc)
plt.plot(fpr_svm, tpr_smv, label='SVM (area = %0.2f)' % svm_roc)
plt.plot(fpr_rand, tpr_rand, linestyle='--', label='Random Baseline')
plt.legend(loc='lower right', prop={'size':18})
plt.show()
```



Which feature is most important?

Extra Tree Classifier

```
In [269]: et = ExtraTreesClassifier(n_estimators=500)
```

```
In [270]: et.fit(X_train, Y_train)
et_test_preds = et.predict_proba(X_test)[:, 1]
et_preds = et.predict(X_test)
```

```
In [271]: print (pd.crosstab(index=Y_test, columns=et_preds, rownames=['True'],
colnames=['Predicted']))
```


Predicted	0	1
True		
0	177	14
1	14	95

```
In [272]: fpr_et, tpr_et, thresholds_et = roc_curve(Y_test, et_test_preds)
```

```
In [273]: et_roc = roc_auc_score(Y_test, et_test_preds)
roc_auc_score(Y_test, et_test_preds)
```

```
Out[273]: 0.95624189442336327
```

```
In [274]: plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %  
0.2f)' % dt_roc)  
plt.plot(fpr_rf, tpr_rf, label='Random Forest 1 (area = %0.2f)' % rf_roc)  
plt.plot(fpr_rf_custom, tpr_rf_custom, label='Random Forest 2 (area = %0.2f)'  
% rf_custom_roc)  
plt.plot(fpr_gbc, tpr_gbc, label='Gradient Boosting (area = %0.2f)' % gbc_roc)  
plt.plot(fpr_svm, tpr_smv, label='SVM (area = %0.2f)' % svm_roc)  
plt.plot(fpr_et, tpr_et, label='Extra Trees (area = %0.2f)' % et_roc)  
plt.plot(fpr_rand, tpr_rand, linestyle='--', label='Random Baseline')  
plt.legend(loc='lower right', prop={'size':18})  
plt.show()
```



```
In [275]: #neural network clasificer
clfnn = MLPClassifier(solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(5, 2), r
andom_state=1)
clfnn.fit(X_train, Y_train)
MLPClassifier(activation='relu', alpha=1e-05, batch_size='auto',
              beta_1=0.9, beta_2=0.999, early_stopping=False,
              epsilon=1e-08, hidden_layer_sizes=(5, 2), learning_rate='constant',
              learning_rate_init=0.001, max_iter=200, momentum=0.9,
              nesterovs_momentum=True, power_t=0.5, random_state=1, shuffle=True,
              solver='lbfgs', tol=0.0001, validation_fraction=0.1, verbose=False,
              warm_start=False)
nn_test_preds = clfnn.predict_proba(X_test)[: , 1]
nn_preds = clfnn.predict(X_test)
print (pd.crosstab(index=Y_test, columns=nn_preds, rownames=['True'],
colnames=['Predicted']))
```

	Predicted	0	1
True			
0		180	11
1		11	98


```
In [276]: nn_roc = roc_auc_score(Y_test, nn_test_preds)
roc_auc_score(Y_test, nn_test_preds)
```

```
Out[276]: 0.95547336567558483
```

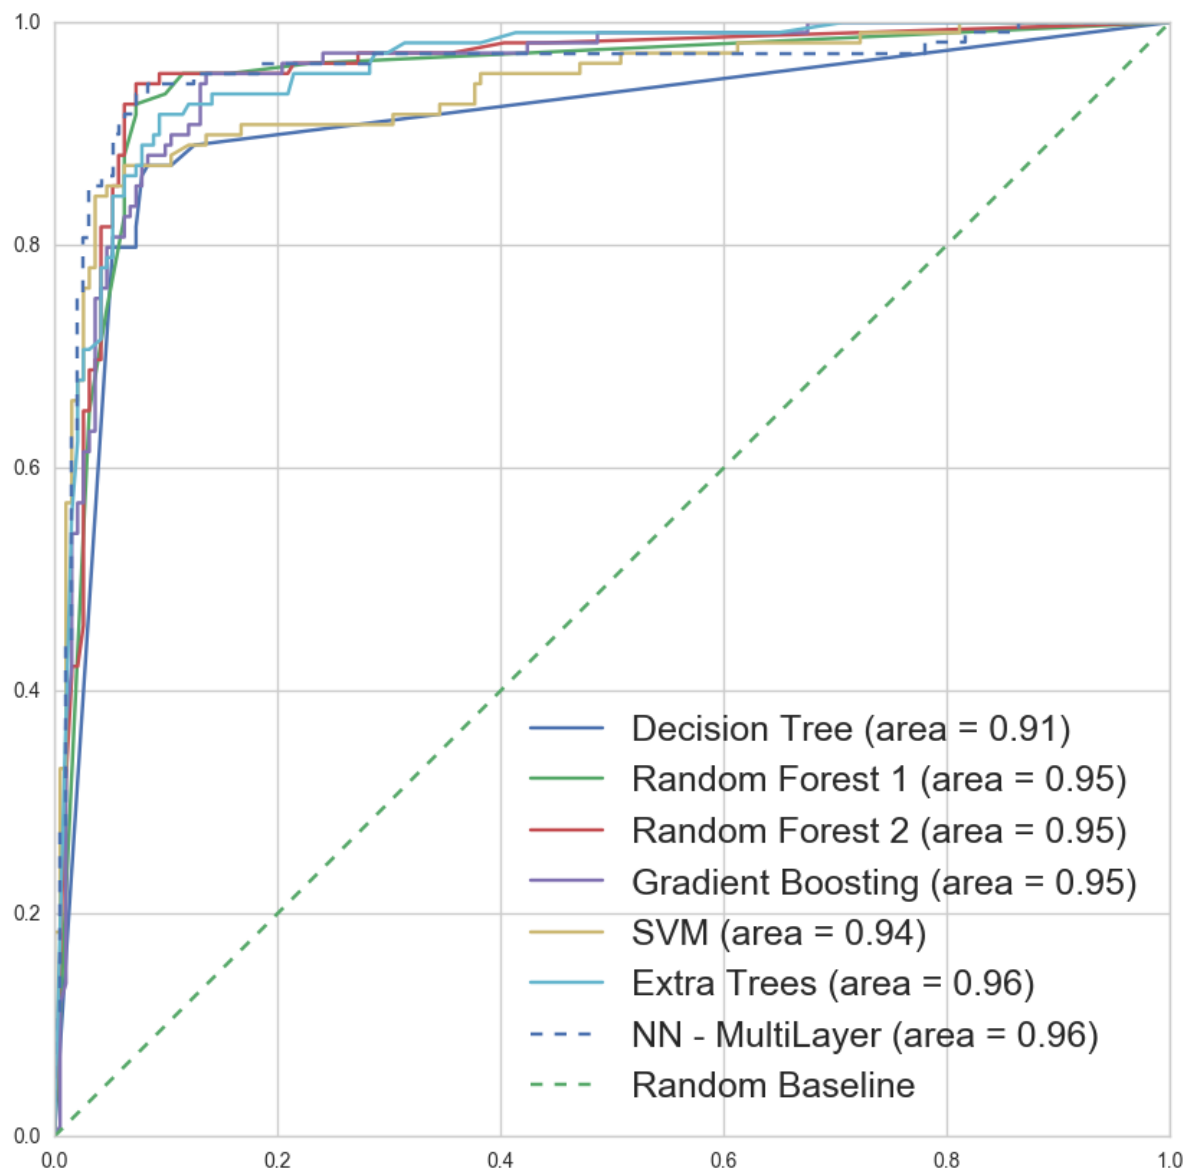


```
In [277]: fpr_nn, tpr_nn, thresholds_nn = roc_curve(Y_test, nn_test_preds)
plt.plot(fpr_nn, tpr_nn, label='NN (area = %0.2f)' % nn_roc)
```

```
Out[277]: [<matplotlib.lines.Line2D at 0x2cab58de1d0>]
```



```
In [278]: plt.plot(fpr_decision_tree, tpr_decision_tree, label='Decision Tree (area = %
0.2f)' % dt_roc)
plt.plot(fpr_rf, tpr_rf, label='Random Forest 1 (area = %0.2f)' % rf_roc)
plt.plot(fpr_rf_custom, tpr_rf_custom, label='Random Forest 2 (area = %0.2f)'
% rf_custom_roc)
plt.plot(fpr_gbc, tpr_gbc, label='Gradient Boosting (area = %0.2f)' % gbc_roc)
plt.plot(fpr_svm, tpr_smv, label='SVM (area = %0.2f)' % svm_roc)
plt.plot(fpr_et, tpr_et, label='Extra Trees (area = %0.2f)' % et_roc)
plt.plot(fpr_nn, tpr_nn, linestyle='--', label='NN - MultiLayer (area = %0.2f)
nn_roc)
plt.plot(fpr_rand, tpr_rand, linestyle='--', label='Random Baseline')
plt.legend(loc='lower right', prop={'size':18})
plt.show()
```



Model ROC Summary

```
In [279]: print("ROC: %0.3f (Decision Tree)" % (dt_roc))
print("ROC: %0.3f (Random Forest - Default Parameters)" % (rf_roc))
print("ROC: %0.3f (Random Forest - Custom Parameters)" % (rf_custom_roc))
print("ROC: %0.3f (Gradient Boosting)" % (gbc_roc))
print("ROC: %0.3f (SVM)" % (svm_roc))
print("ROC: %0.3f (Extra Trees)" % (et_roc))
print("ROC: %0.3f (NN MultiLayer)" % (nn_roc))
```

```
ROC: 0.909 (Decision Tree)
ROC: 0.948 (Random Forest - Default Parameters)
ROC: 0.955 (Random Forest - Custom Parameters)
ROC: 0.952 (Gradient Boosting)
ROC: 0.938 (SVM)
ROC: 0.956 (Extra Trees)
ROC: 0.955 (NN MultiLayer)
```

Based on the above ROC, The Extra Trees followed by MultiLayer Neural Network are the preferred model to predict the Severity of incidents